

Log Management in Apache Kafka

Lee Dongjin | dongjin@apache.org

Introducing Myself ...

- Software Developer, Big Data Technologies
 - Hadoop, Giraph, HBase, Spark, Kafka, ...
- Apache Kafka Consultant
- Apache Kafka Contributor
 - [KIP-110: Add Codec for ZStandard Compression](#)
 - [KIP-508: Make Suppression State Queriable](#)
 - [KIP-653: Upgrade log4j to log4j2 \(Today's Topic!\)](#)
 - etc...

Log - What Log?

- Log = Set of events
- Meaning 1:
 - Append-only data structure
 - In Kafka, stored in `${log.dirs}`, `${log.dir}`
- Meaning 2:
 - A file that records the events occur in the system internals.
 - In Kafka, stored in `${kafka.logs.dir}`
- We will focus on meaning 2 only in this talk.

Log Management in Kafka

- Logging in Kafka Broker
 - What to monitor?
 - How to manage it?
- Logging in Kafka Connect
 - What to monitor?
 - How to manage it?
- Upcoming: Log4j2 (Preview)
 - Differences?
 - How to activate it?

Logging in Kafka Broker (1)

- So many loggers in Kafka Broker
 - {number of loggers} > 110 (as of 2.7.0)
- Too many loggers to keeping eyes on them...
 - Then, which loggers are more important?

Logging in Kafka Broker (2)

- `kafka.request.logger`
 - Receives various Kafka requests.
 - Default: `${kafka.logs.dir}/kafka-request.log`
- `kafka.server.KafkaApis`
 - Handles various Kafka requests.
 - Not defined in the default logging configuration.
- `kafka.log.LogCleaner`
 - Log Compaction, Retention, etc.
 - Default: `${kafka.logs.dir}/log-cleaner.log`

How to manage logging in Kafka Broker? (1)

- Modifying the log configuration & restarting?
 - A squad of brokers?
 - Kubernetes pods?
- We need to configure the log settings dynamically!
- Alternatives
 - Method 1: Use JMX Console
 - Method 2: Use Command Line Tool (kafka-configs.sh)
 - ... or Admin API

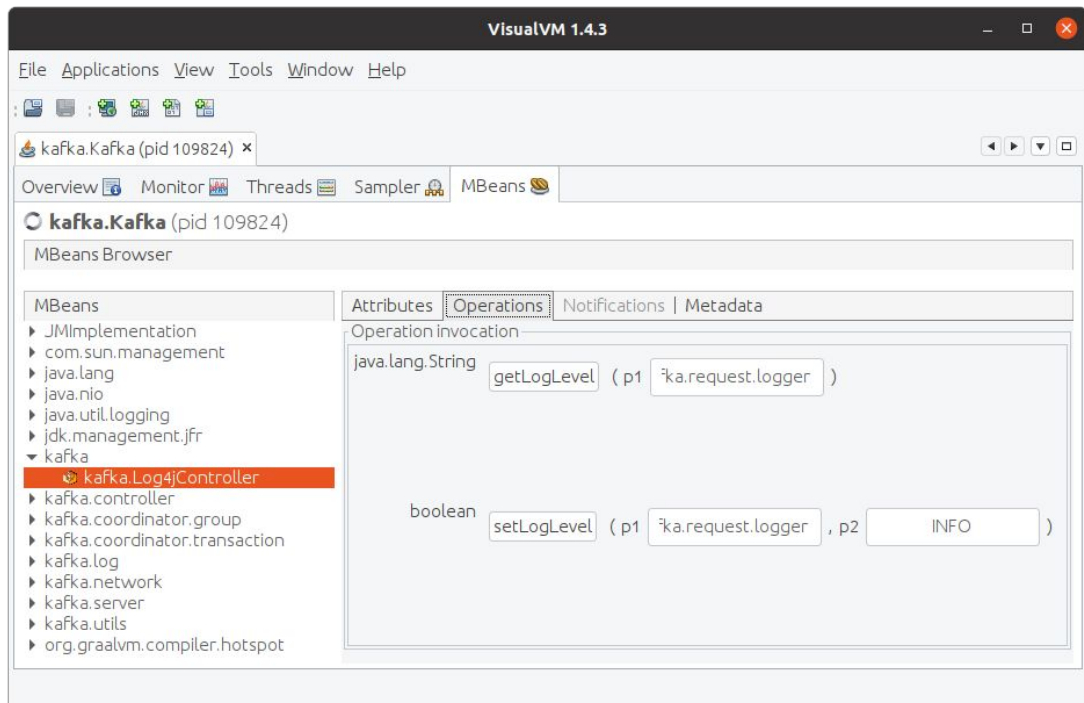
How to manage logging in Kafka Broker? (2)

- Method 1: Use JMX Console
 - You can change the logger's logging level dynamically with JMX bean.
 - Similar Features: Apache Storm, Apache Zookeeper
- How to use it?

How to manage logging in Kafka Broker? (3)

- Start Kafka with JMX enabled.
- Connect to {kafka-broker}:{jmx-port} with JMX Console
- Open `kafka > kafka.request.logger` MBean
 - To get Logger Level
 - click `getLogLevel` from operations tab
 - `p1` = logger name (e.g., 'kafka.request.logger') ↵
 - To set Logger Level
 - click `setLogLevel` from operations tab
 - `p1` = logger name (e.g., 'kafka.request.logger'), `p2` = logger level (e.g., 'INFO') ↵

How to manage logging in Kafka Broker? (4)



How to manage logging in Kafka Broker? (5)

- Method 2: Use Command Line Tool
 - To check current loggers:

```
$ bin/kafka-configs.sh --bootstrap-server {kafka-broker}:9092 \  
    --entity-type broker-loggers --entity-name {broker-id} \  
    --describe
```

Dynamic configs for broker-logger {broker-id} are:

```
org.apache.zookeeper.CreateMode=INFO sensitive=false synonyms={}  
kafka.coordinator.transaction.ProducerIdManager=INFO sensitive=false synonyms={}  
kafka.utils.KafkaScheduler=INFO sensitive=false synonyms={}  
...
```

How to manage logging in Kafka Broker? (6)

- Method 2: Use Command Line Tool
 - To update current logger level (e.g., `kafka.request.logger` to `TRACE`):

```
$ bin/kafka-configs.sh --bootstrap-server {kafka-broker}:9092 \  
    --entity-type broker-loggers --entity-name {broker-id} \  
    --alter --add-config kafka.request.logger=TRACE  
Completed updating config for broker-logger {broker-id}.
```

How to manage logging in Kafka Broker? (7)

- Dynamic Configuration API was updated in 2.3.0
 - Admin#incrementalAlterConfigs
 - [KIP-339: Create a new IncrementalAlterConfigs API](#)
 - [KIP-412: Extend Admin API to support dynamic application log levels](#)
- Before use this tool, please update the installation first.

Logging in Kafka Connect (1)

- Question: Too many logging messages are mixed in one log file.
 - `${kafka.logs.dir}/connect.log`
 - How can I separate the logging messages I need...?
 - Like, the log messages from a specific task.
- KIP-449: Add connector contexts to Connect worker logs
 - From 2.3.0
 - Provides Logging Context

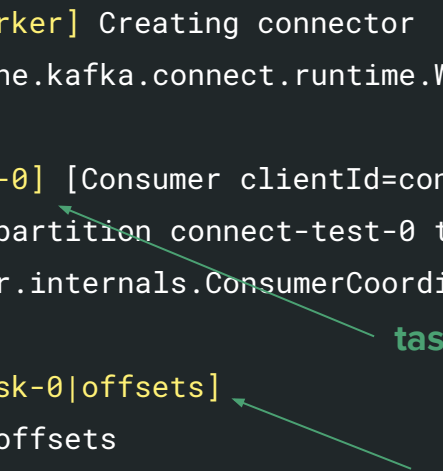
Logging in Kafka Connect (2)

- Logging format: [<connectorName>|<scope>] <message>
 - [my-jdbc-connector|worker] :
 - Connector Creation, Config Validation, Starting, Stopping, ...
 - [my-jdbc-connector|task-0]:
 - Task Creation, Consuming, Producing, ...
 - [my-jdbc-connector|task-0|offsets]:
 - Task's offset commit.

Logging in Kafka Connect (3)

- Example

```
[2019-04-02 17:01:38,315] INFO [local-file-source|worker] Creating connector
local-file-source of type FileStreamSource (org.apache.kafka.connect.runtime.Worker:227)
...
[2019-04-02 17:01:38,487] INFO [local-file-sink|task-0] [Consumer clientId=consumer-1,
groupId=connect-local-file-sink] Setting offset for partition connect-test-0 to the
committed offset 2 (org.apache.kafka.clients.consumer.internals.ConsumerCoordinator:511)
...
[2019-04-02 17:01:40,901] INFO [local-file-source|task-0|offsets]
WorkerSourceTask{id=local-file-source-0} Committing offsets
(org.apache.kafka.connect.runtime.WorkerSourceTask:398)
```



The diagram consists of three green arrows pointing from text labels to specific parts of the log output. The first arrow, labeled 'worker log', points to the log entry '[2019-04-02 17:01:38,315] INFO [local-file-source|worker]'. The second arrow, labeled 'task log', points to the log entry '[2019-04-02 17:01:38,487] INFO [local-file-sink|task-0]'. The third arrow, labeled 'task commits offset!', points to the log entry '[2019-04-02 17:01:40,901] INFO [local-file-source|task-0|offsets]'.

How to manage logging in Kafka Connect? (1)

- Checking current loggers with corresponding levels
 - GET to `http://{kafka-connect-host}:8083/admin/loggers`
- Example

```
$ curl -s http://{kafka-connect-host}:8083/admin/loggers
{
  "org.apache.kafka.connect.runtime.rest": {
    "level": "WARN"
  },
  ...
}
```

How to manage logging in Kafka Connect? (2)

- Modifying Logging Levels
 - Put Json to `http://{kafka-connect-host}:8083/admin/loggers/{logger-name}`
- Example
 - Change logger `'io.debezium.*'`s level into `'TRACE'`

```
$ curl -s -X PUT -H "Content-Type:application/json" \  
    http://{kafka-connect-host}:8083/admin/loggers/io.debezium \  
    -d '{"level": "TRACE"}' | jq '.'  
  
[  
  "io.debezium",  
  "io.debezium.connector.mysql.MySqlConnector",  
  ...  
]
```

Upcoming Feature: Log4j2 (1)

- Log4j 1.x: deprecated from May, 2012 (1.2.17)
 - Obsolete Configuration Syntax
 - Security Problems
 - [CVE-2019-17571: SocketServer Vulnerability](#)
 - **Problem: Kafka is still using it.**
 - Only Log4j 1.x to Kafka is supported.
- Log4j 2.x: current de-facto standard
 - Will be adopted to Apache Kafka > 2.7.0
 - [KIP-676: Respect logging hierarchy](#) (Bug Fix)
 - [KIP-653: Upgrade log4j to log4j2](#)

Upcoming Feature: Log4j2 (2)

- log4j 1.x vs. log4j2 Configuration Syntax

```
# root logger
log4j.rootLogger=INFO, stdout, kafkaAppender

...

log4j.logger.kafka=INFO
log4j.logger.org.apache.kafka=INFO

log4j.logger.kafka.request.logger=WARN,
requestAppender
log4j.additivity.kafka.request.logger=false
```

```
# root logger
rootLogger.level=INFO

...

# in log4j2, list loggers first
loggers=org.apache.zookeeper, ...

logger.org.apache.kafka.name=org.apache.kafka
logger.org.apache.kafka.level=INFO
logger.kafka.request.logger.additivity=false
```

Upcoming Feature: Log4j2 (3)

- Replaces log4 to log4j2
 - Removes Security Vulnerability
- Default: Log4j configuration
 - For backward-compatibility (with log4j-1.2-api)
- How can I activate log4j2 mode?
 - Use JVM configuration
 - Log4j: log4j.configuration
 - Log4j2: log4j.configuration**File**
 - export

KAFKA_LOG4J_OPTS="-Dlog4j.configurationFile=file:\$base_dir/../../config/log4j2.properties"

Upcoming Feature: Log4j2 (3)

- What Changes?
 - Except Configuration Syntax, NONE.
 - Log4j2 version default configurations will be provided, also.
- Note: Root Logger's name
 - Log4j: "root"
 - Log4j2: "" (empty string!)
 - Kafka: "root"
 - For backward-compatibility

Upcoming Feature: Log4j2 (3)

- Question 1: When can we use the feature?
 - Public Beta will be provided soon.
 - 2.7.0+log4j2, 2.6.1+log4j2, etc.
 - with English, Korean documentation.
- Question 2: When log4j2 configuration will be default?
 - Not Decided Yet.
- Question 3: What about Log4j2 Appender?
 - Working in Progress.

Summary

- You can manage Kafka Brokers' log with ...
 - JMX Console
 - Command Line Tool
- You can manager Kafka Connects' log with...
 - Rest API
- Kafka will migrate into log4j2
 - With Backward-compatibility!

Questions?