

# Kafka 사용사례

2018. 09. 08.

류지형  
Kaiser.Ryu



# 목차


- ✓ Who am I
- ✓ 모니터링 시스템
- ✓ 서비스 구성도
- ✓ System 구성
- ✓ Dashborad
- ✓ Future work

# **PART 1.**

## **Who am I**

# Who am I

- 류지형(Kaiser.Ryu)
- SK텔레콤 Datalake 운영
- JAVA, Hadoop, Hive, HBase  
MySQL, Bash shell, Python
- Kafka 경험 약 6개월




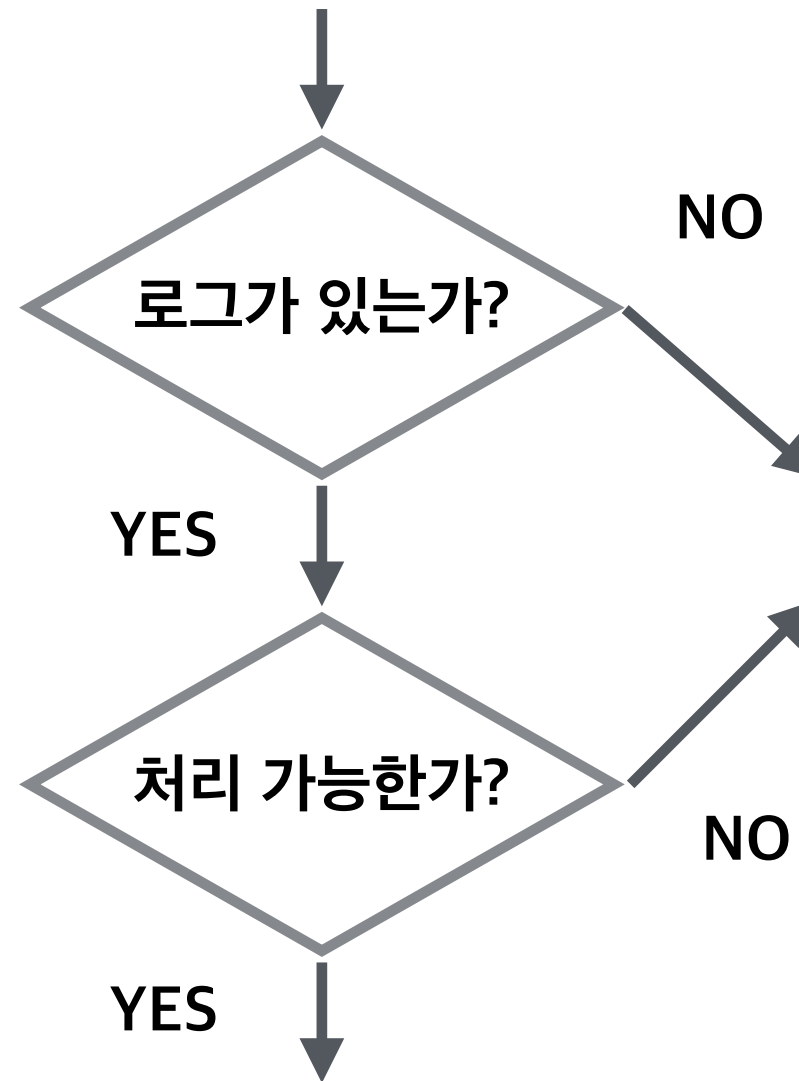
# **PART 2.**

## **모니터링 시스템**

# 모니터링 시스템

## 왜 필요한가?

- 장애예방
- 빠른 장애 인지
- 정확한 원인 파악
- 신속한 처리
- 재발 방지



**로그 및 메트릭 수집은 매우 중요**

# **PART 3.**

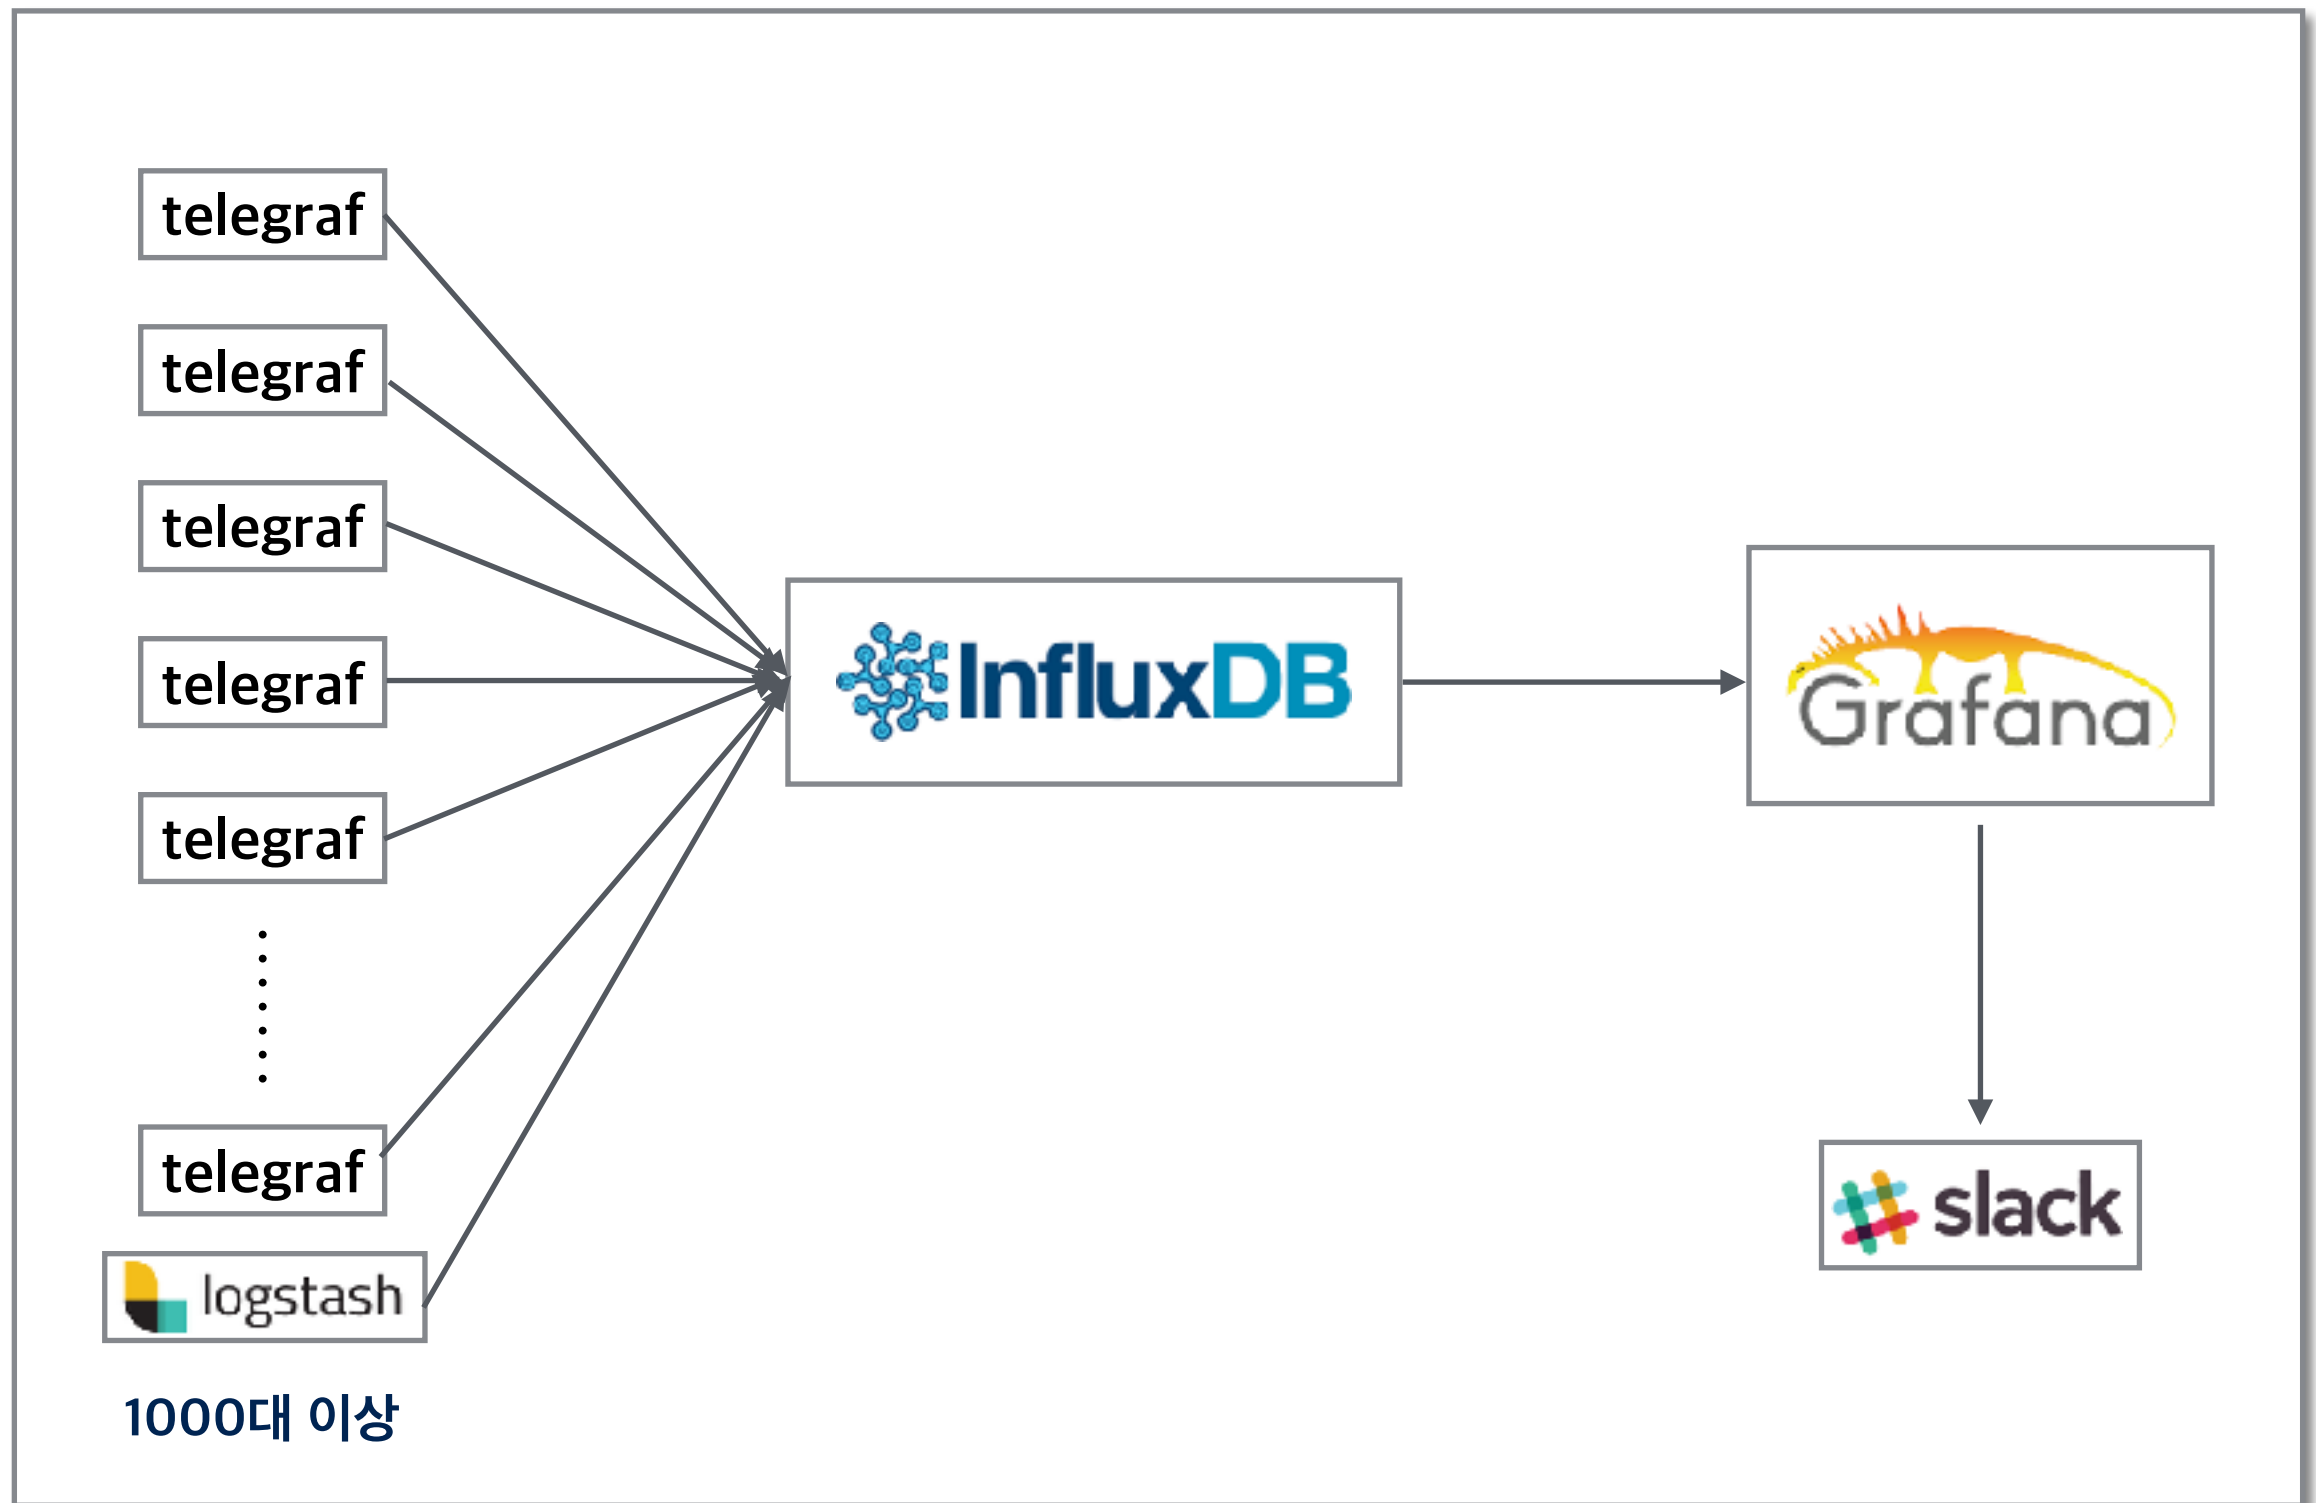
## **서비스 구성도**



# AS-IS

수집 서버 1대가 1000대 이상 서버의 metric을 전송받는 형태

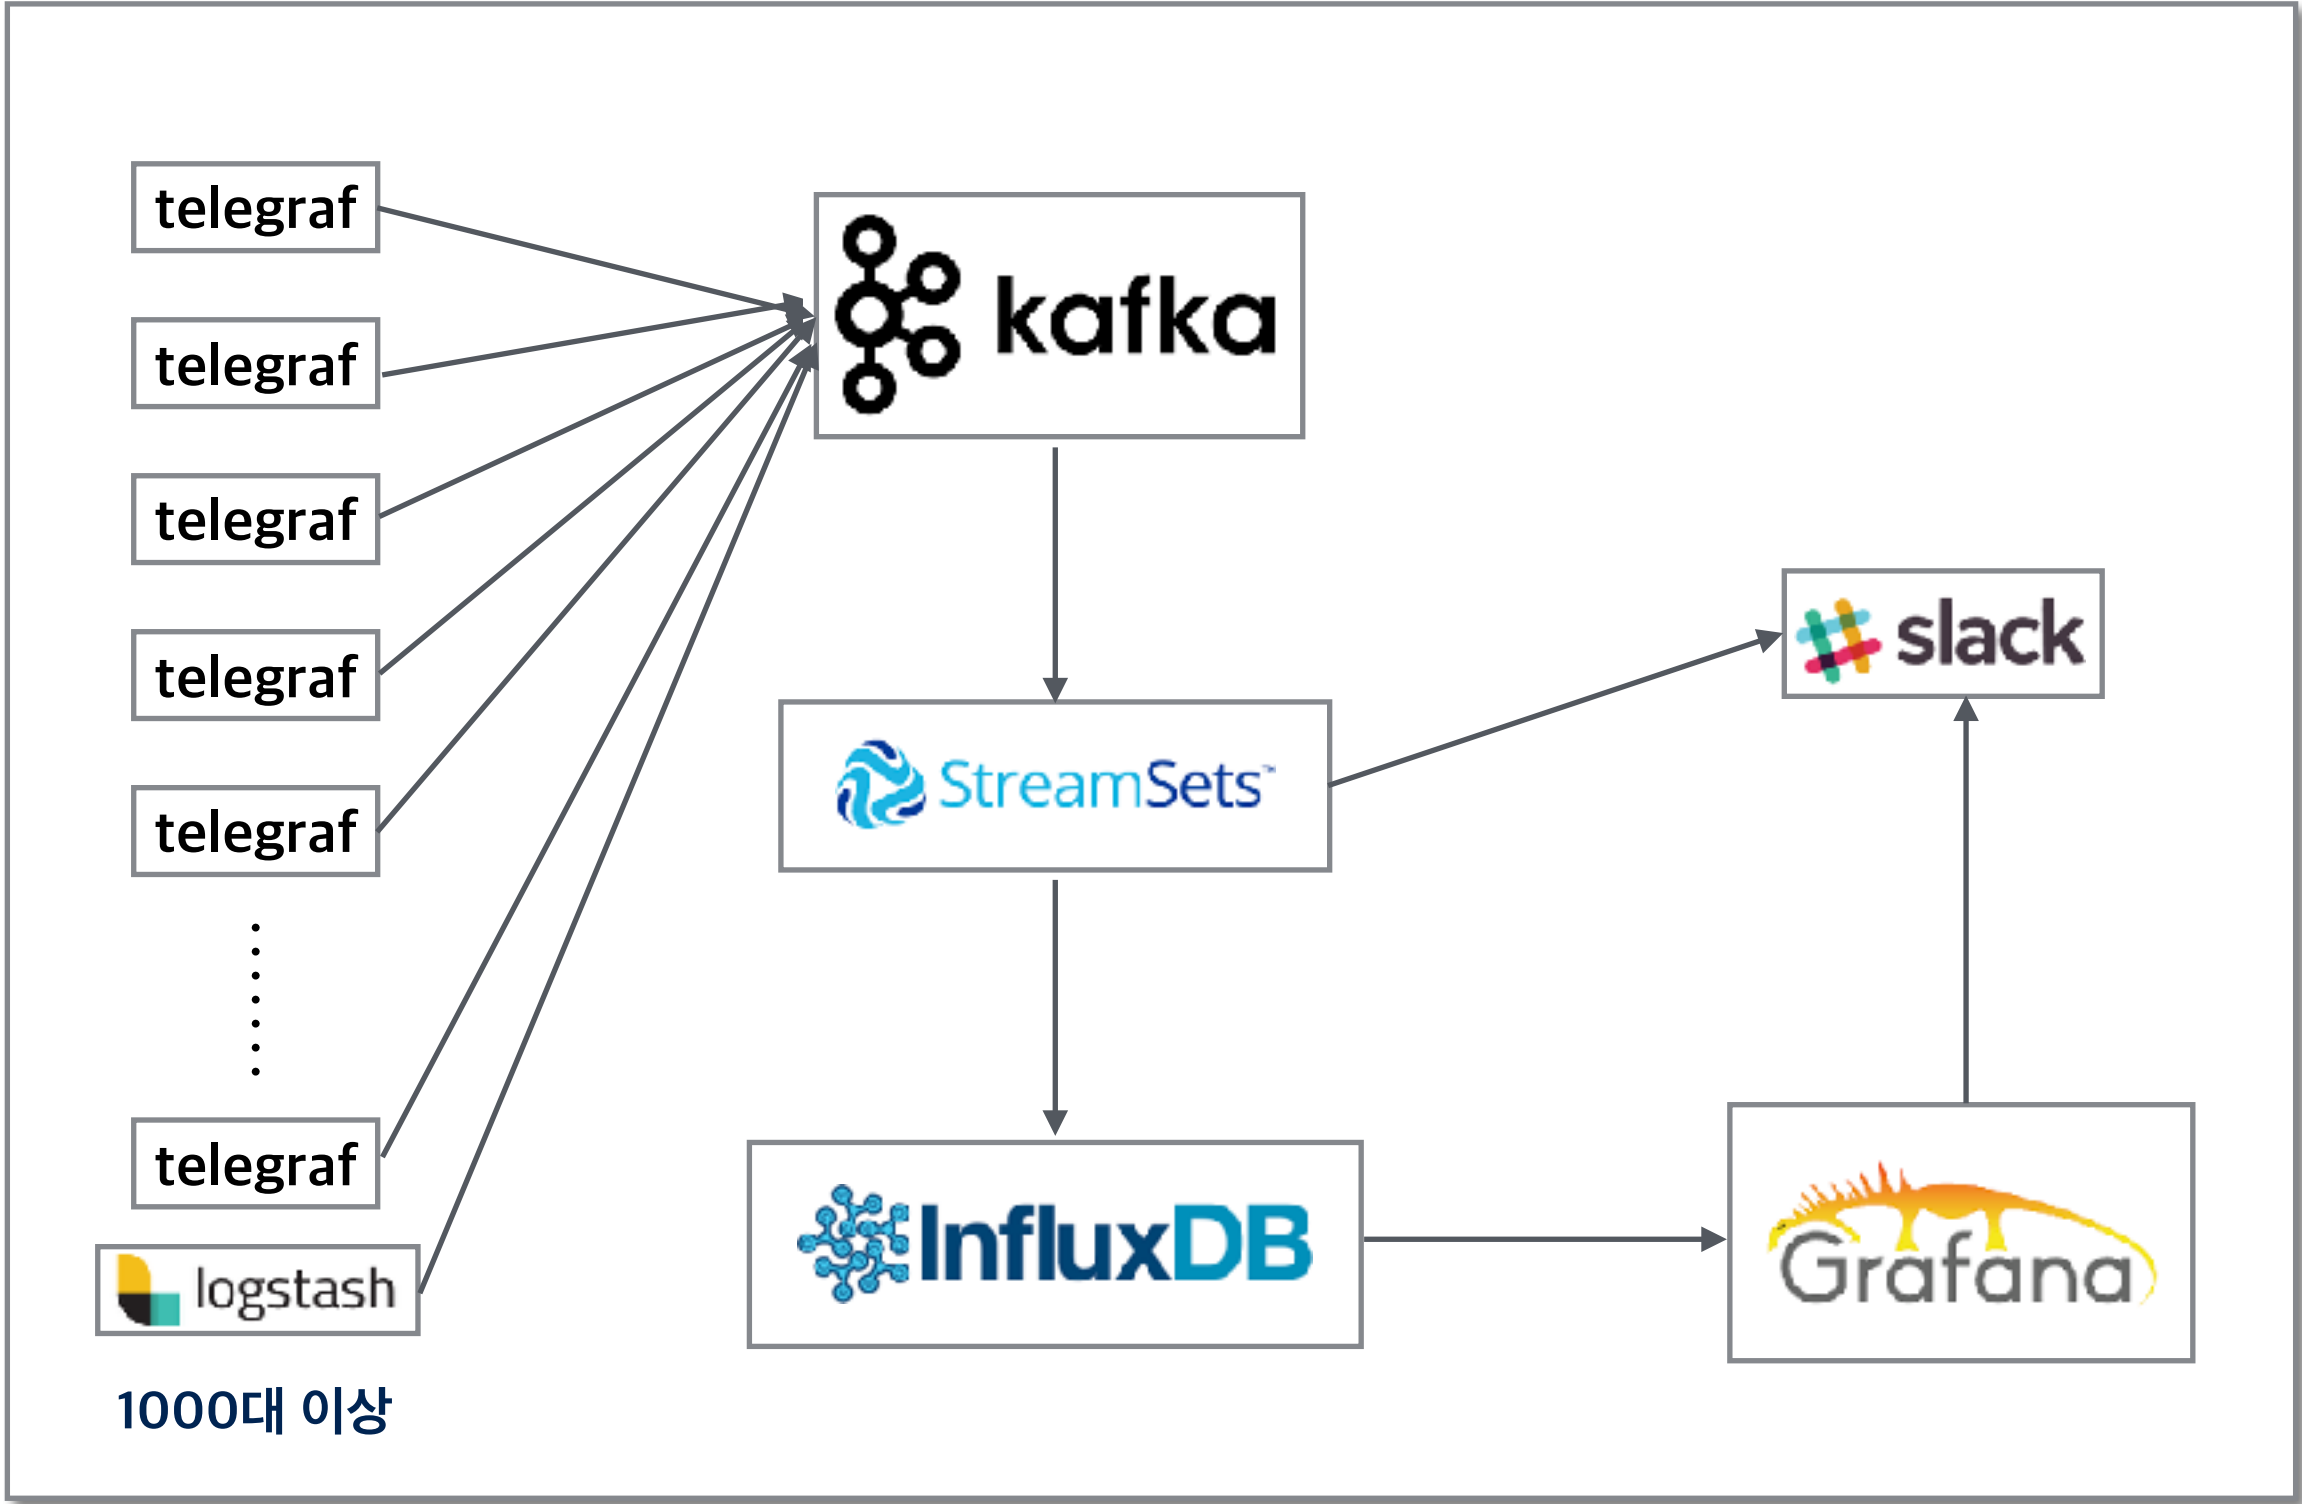
수집 서버가 다운될 경우 해당 시간동안 모든 metric 유실



# TO-BE

모든 클러스터의 mertric을 카프카로 전송.

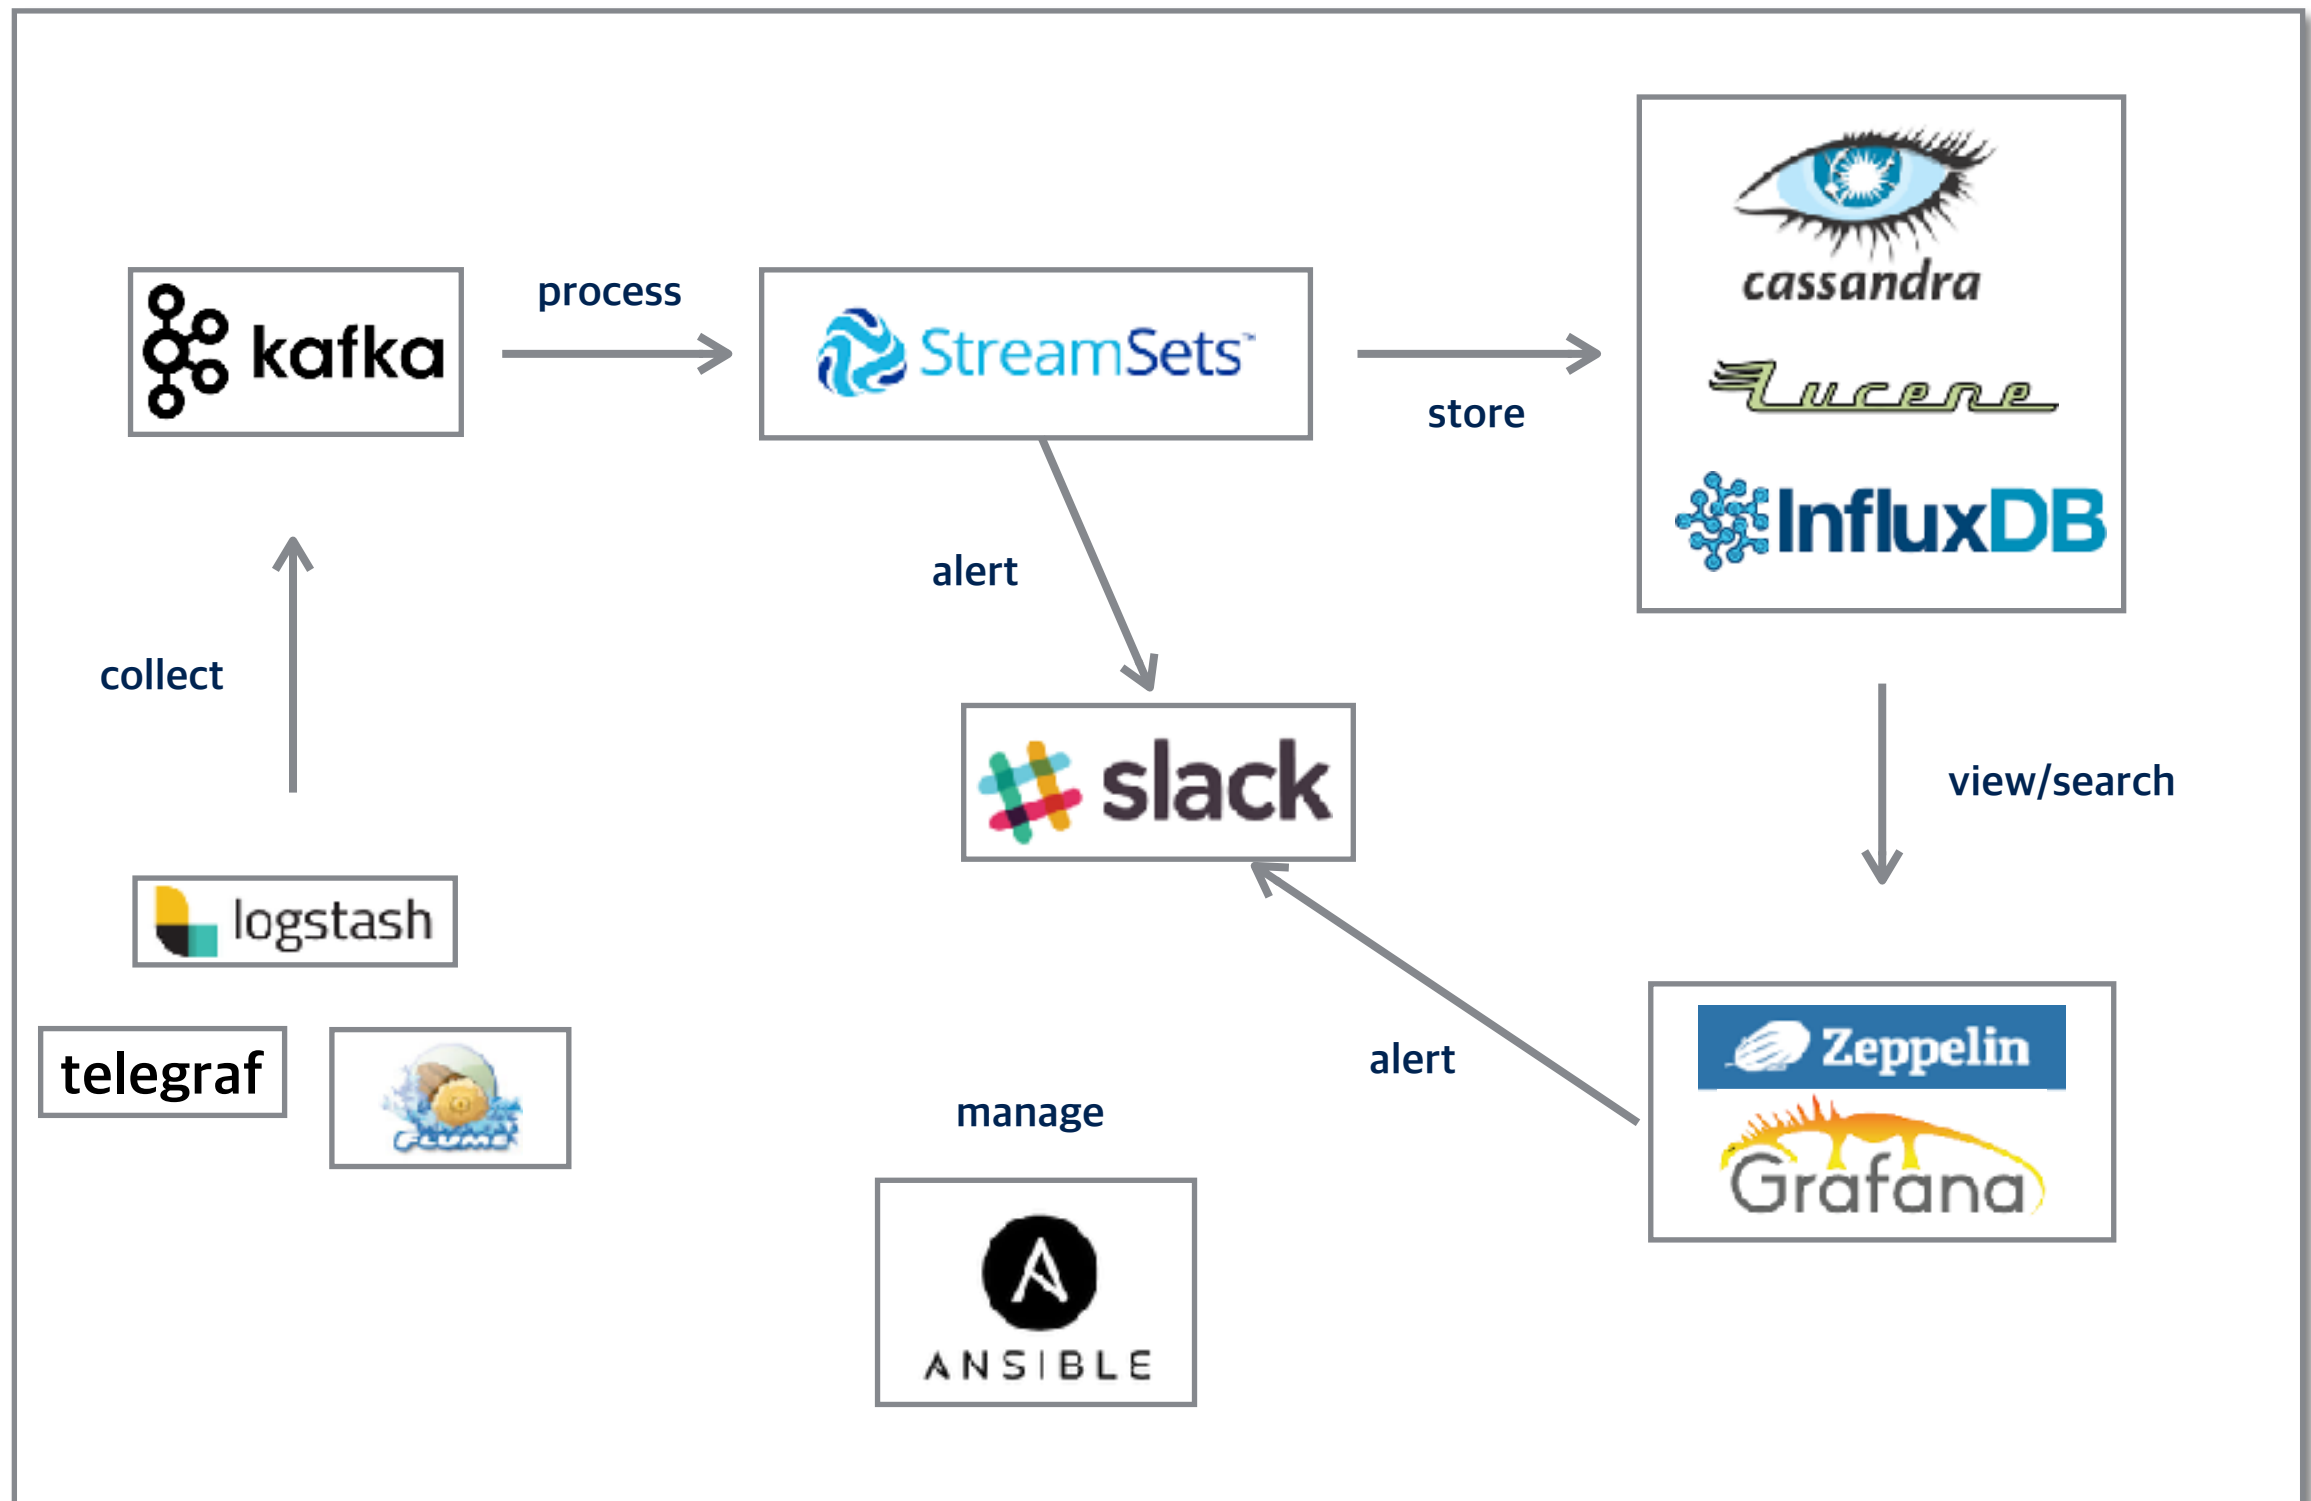
SPOF를 줄이고 필요한 경우 StreamSets을 이용하여 전처리 및 Alert 전송



# 서비스 구성도

telegraf, logstash, flume을 이용하여 metric 및 로그 kafka 전송

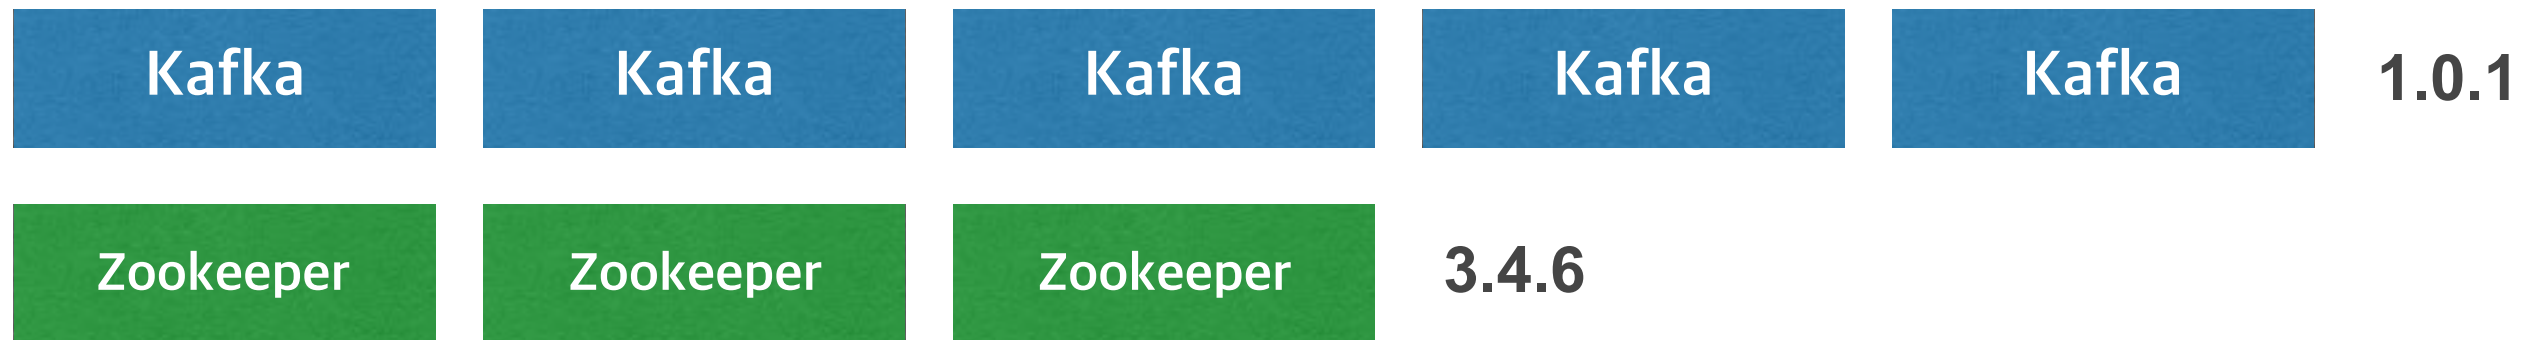
influxDB 및 cassandra에 적재된 데이터를 grafana, zeppelin을 이용하여 시각화



# PART 4.

## System 구성

# Kafka Broker



- Disk 2TB x 2 (총 20TB 구성)
- Topic 1개로 구성
  - cpu, memory, network, disk, hadoop metric, jvm metric, switch metric ...
- Replication Factor 3
- Partition 5
- Retention 7days

6TB Storage 사용  
850GB per day  
35GB per hour

( 3 factor )

# Kafka Producer(telegraf)

## telegraf.conf

```
[[outputs.kafka]]
```

```
## URLs of kafka brokers
```

```
brokers = ["kafka-01-01:9092","kafka-01-02:9092"]
```

```
## Kafka topic for producer messages
```

```
topic = "datalake"
```

```
## CompressionCodec represents the various compression codecs recognized by  
## Kafka in messages.
```

```
## 0 : No compression
```

```
## 1 : Gzip compression
```

```
## 2 : Snappy compression
```

```
compression_codec = 0
```

```
## RequiredAcks is used in Produce Requests to tell the broker how many  
## replica acknowledgements it must see before responding
```

```
## 0 : the producer never waits for an acknowledgement from the broker.
```

```
## This option provides the lowest latency but the weakest durability  
## guarantees (some data will be lost when a server fails).
```

```
## 1 : the producer gets an acknowledgement after the leader replica has  
## received the data. This option provides better durability as the  
## client waits until the server acknowledges the request as successful  
## (only messages that were written to the now-dead leader but not yet  
## replicated will be lost).
```

```
## -1: the producer gets an acknowledgement after all in-sync replicas have  
## received the data. This option provides the best durability, we  
## guarantee that no messages will be lost as long as at least one in  
## sync replica remains.
```

```
required_acks = 1
```

```
## The total number of times to retry sending a message
```

```
max_retry = 3
```

```
## https://github.com/influxdata/telegraf/blob/master/docs/DATA\_FORMATS\_OUTPUT.md
```

```
data_format = "influx"
```

# Kafka Producer(Logstash)

## switch\_metric.conf

```
event['now'] = Time.now.to_i * 1000000000;
```

```
output {  
  kafka {  
    topic_id => "datalake"  
    bootstrap_servers => "kafka-01-01:9092,kafka-01-02:9092"  
    acks => "1"  
    codec => plain {  
      format => "switch,inout=%{inout} 1001=%{1001},2001=%{2001} %{now}"  
    }  
  }  
}
```

## influx 포맷

```
switch,inout=in 1001=0i,2001=0i 1525417999000000000  
switch,inout=out 1001=8848634164655i,2001=8373785484971i 1525417999000000000  
switch,inout=in 1001=8625697583336i,2001=9182721564736i 1525417999000000000  
switch,inout=out 1001=0i,2001=0i 1525417999000000000
```

# Kafka Consumer(StreamSets)

Kafka Consume => Record Merge => InfluxDB Insert 의 과정으로 수행된다.

(Record Merge는 1000개 Row 기준)

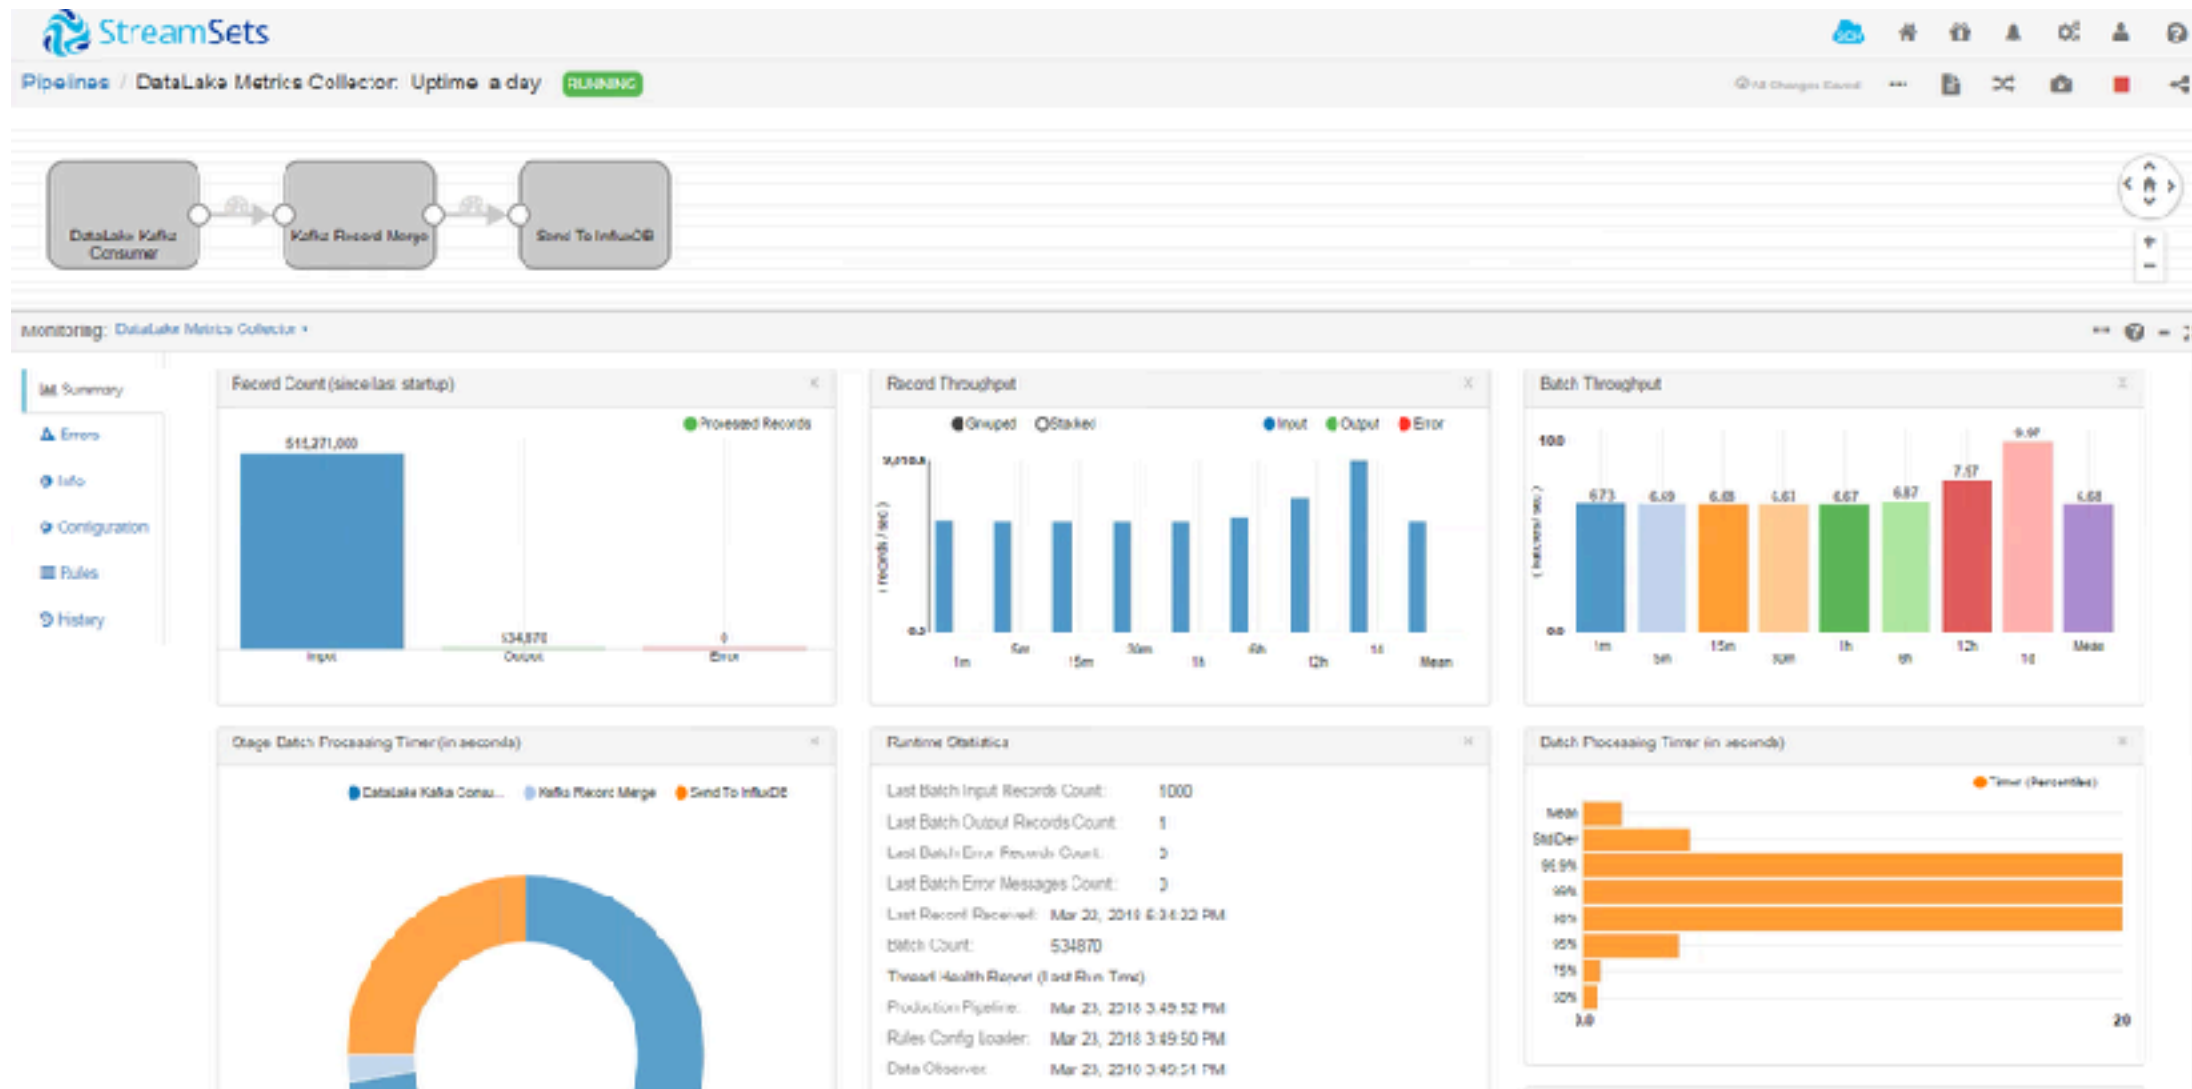
## Configuration

Broker URI : kafka-01-01:9092,kafka-01-02:9092

ZooKeeper URI : kafka-01-03:2181

Consumer Group : datalakeConsumerGroup

Topic : datalake





# Kafka Management(Ansible)

## ansible을 이용하여 kafka 기동 스크립트 작성

### start, status, stop 스크립트

#### 1.start\_kafka\_nodes.yml

- **name**: kafka broker nodes
- hosts**: kafka\_nodes

**tasks:**

- **name**: kafka node start
- command**: su kafka -c "/data/kafka/bin/kafka-server-start.sh -daemon /data/kafka/config/server.properties"

#### 2.status\_kafka\_nodes.yml

- **name**: kafka broker nodes
- hosts**: kafka\_nodes

**tasks:**

- **name**: kafka node status
- command**: /data/kafka/bin/kafka-server-status.sh
- register**: status
- **debug**: msg="{{ status.stdout }}"

#### 3.stop\_kafka\_nodes.yml

- **name**: kafka broker nodes
- hosts**: kafka\_nodes

**tasks:**





- **name**: kafka node stop
- command**: /data/kafka/bin/kafka-server-stop.sh

# **PART 5.**

## **Dashboard**



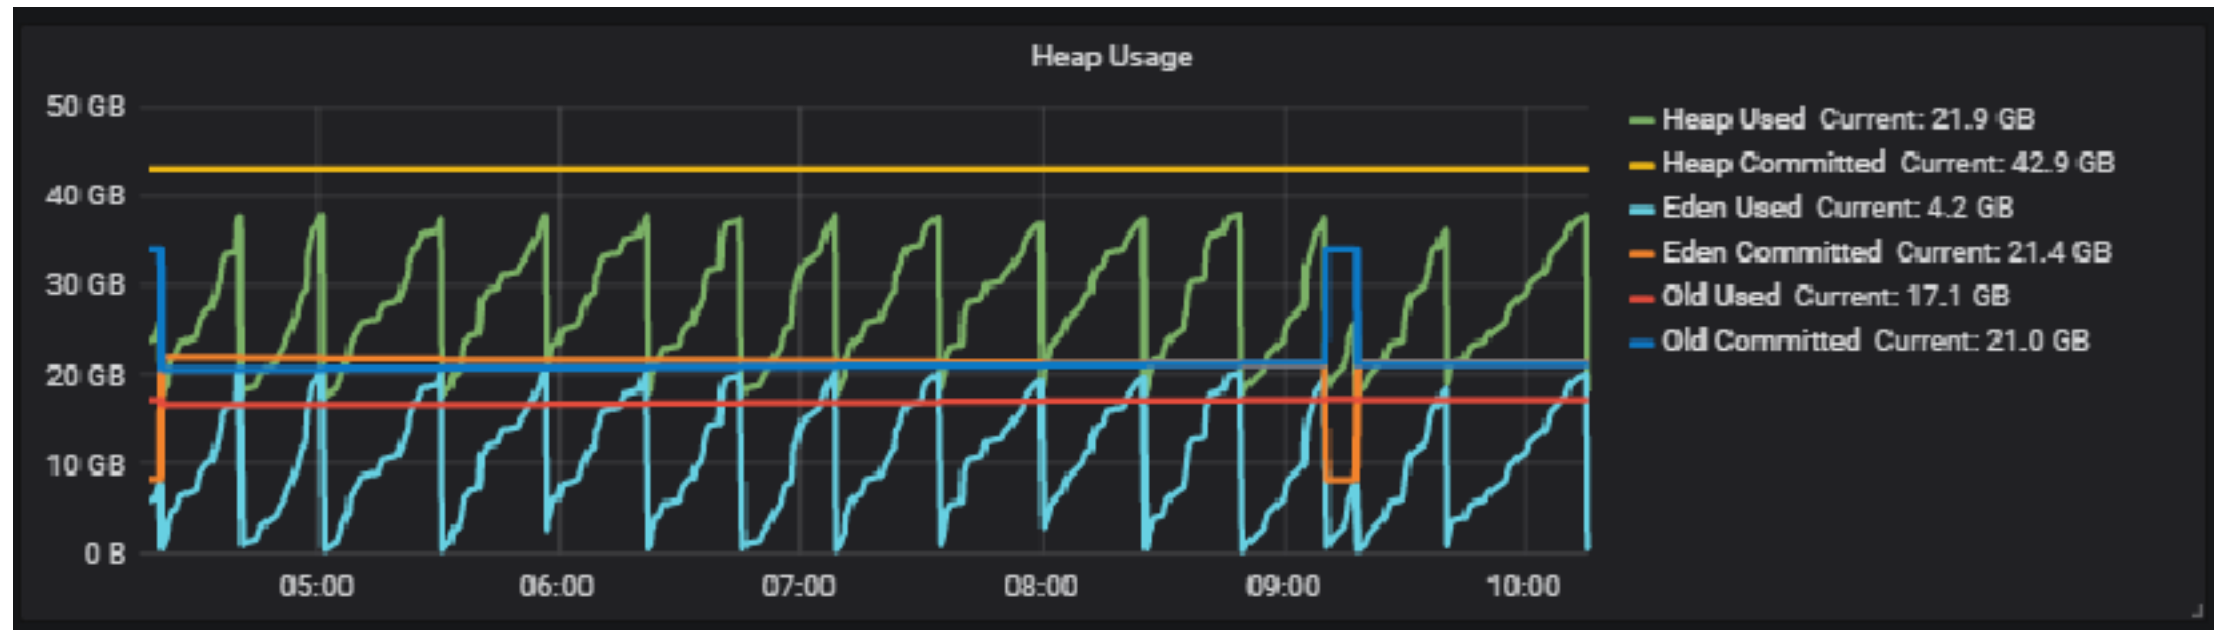
# Grafana

## Hadoop 메트릭



# Grafana

## JVM 메트릭



# **PART 6.**

## **Future work**

그런데 카프카 모니터링은?

실시간 처리에서의 카프카는?

**Thank You!!!**  
**Q & A**